

LARGE SCALE DATA PROCESSING – DIGITAL ASSIGNMENT 2

CASE STUDY: TWITTER DATA STREAMING USING APACHE FLUME

[BEFORE the Twitter data streaming implementation, we understand the concepts of Flume]

Flume: Introduction

Apache Flume is a service for streaming logs into Hadoop. Flume is distributed, reliable, and available service for efficiently collecting, aggregating, and moving large amounts of streaming data into the Hadoop Distributed File System (HDFS). It has a simple and flexible architecture based on streaming data flows; and is robust and fault tolerant with tuneable reliability mechanisms for failover and recovery. YARN coordinates data ingest from Apache Flume and other services that deliver raw data into a Hadoop cluster.

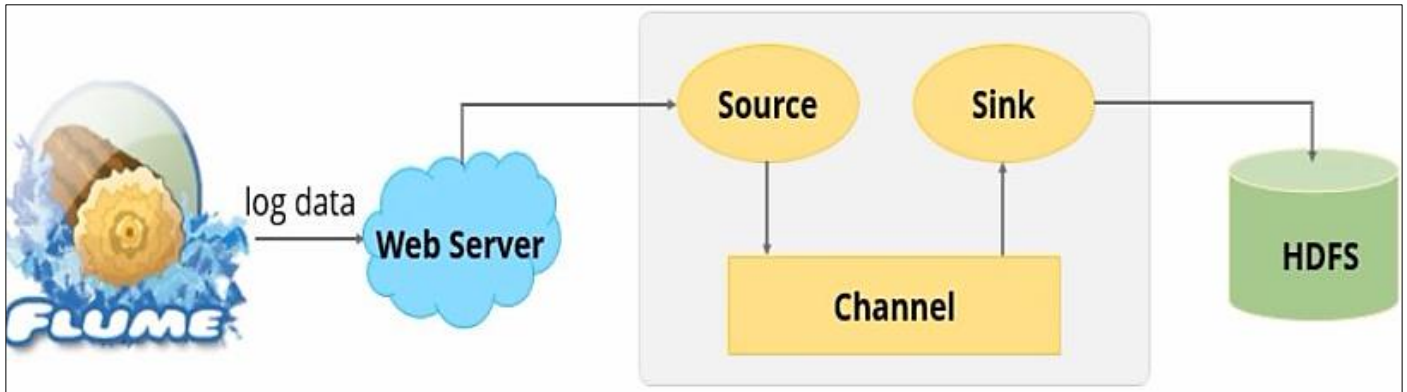
Flume: Features

Flume lets Hadoop users ingest high-volume streaming data into HDFS for storage. Specifically, Flume allows users to:

FEATURES	DESCRIPTION
Stream data	Ingest streaming data from multiple sources into Hadoop for storage and analysis.
Insulate systems	Buffer storage platform from transient spikes, when the rate of incoming data exceeds the rate at which data can be written to the destination.
Guarantee data delivery	Flume NG uses channel-based transactions to guarantee reliable message delivery. When a message moves from one agent to another, two transactions are started, one on the agent that delivers the event and the other on the agent that receives the event. This ensures guaranteed delivery semantics.
Scale horizontally	To ingest new data streams and additional volume as needed.

Enterprises use Flume's powerful streaming capabilities to land data from high-throughput streams in the HDFS. Typical sources of these streams are application logs, sensor and machine data, geo-location data and social media. These different types of data can be landed in Hadoop for future analysis using interactive queries in Apache Hive. Or they can feed business dashboards served ongoing data by Apache HBase.

Example usecase: Flume is used to log manufacturing operations. When one run of product comes off the line, it generates a log file about that run. Even if this occurs hundreds or thousands of times per day, the large volume log file data can stream through Flume into a tool for same-day analysis with Apache Storm or months or years of production runs can be stored in HDFS and analysed by a quality assurance engineer using Apache Hive.



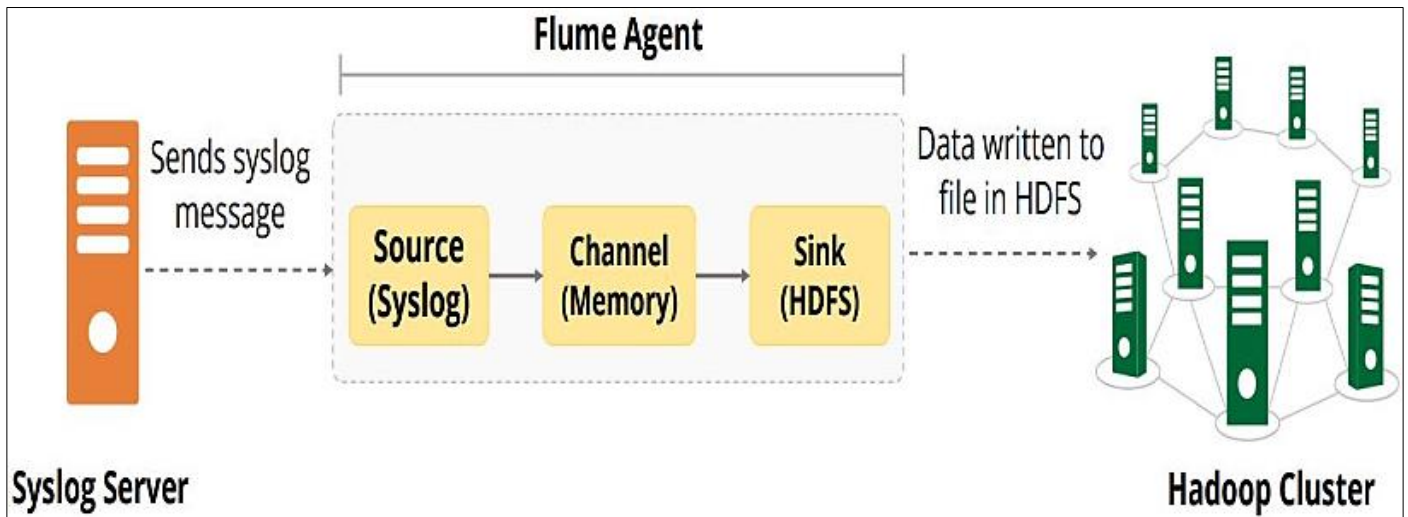
Flume: Working and operations

Flume's high-level architecture is built on a streamlined codebase that is easy to use and extend. The project is highly reliable, without the risk of data loss. Flume also supports dynamic reconfiguration without the need for a restart, which reduces downtime for its agents. The following components make up Apache Flume:

COMPONENT	DEFINITION
Event	A singular unit of data that is transported by Flume (typically a single log entry).
Source	The entity through which data enters into Flume. Sources either actively poll for data or passively wait for data to be delivered to them. A variety of sources allow data to be collected, such as log4j logs and syslogs.
Sink	The entity that delivers the data to the destination. A variety of sinks allow data to be streamed to a range of destinations. One example is the HDFS sink that writes events to HDFS.
Channel	The conduit between the Source and the Sink. Sources ingest events into the channel and the sinks drain the channel.
Agent	Any physical Java virtual machine running Flume. It is a collection of sources, sinks and channels.
Client	The entity that produces and transmits the Event to the Source operating within the Agent.

Flume components interact in the following way:

1. A flow in Flume starts from the **Client**.
2. The **Client** transmits the **Event** to a **Source** operating within the **Agent**.
3. The **Source** receiving this **Event** then delivers it to one or more **Channels**.
4. One or more **Sinks** operating within the same **Agent** drains these **Channels**.
5. **Channels** decouple the ingestion rate from drain rate using the familiar producer-consumer model of data exchange.
6. When spikes in client side activity cause data to be generated faster than can be handled by the provisioned destination capacity can handle, the **Channel** size increases. This allows sources to continue normal operation for the duration of the spike.
7. The **Sink** of one **Agent** can be chained to the **Source** of another **Agent**. This chaining enables the creation of complex data flow topologies.



How is Flume easily scalable?

Because Flume's distributed architecture requires no central coordination point. Each agent runs independently of others with no inherent single point of failure, and Flume can easily scale horizontally.

FLUME CASE STUDY: TWITTER DATA STREAMING

STEP 1:

Setting up Apache Flume: We download and install **Flume 1.9**. We also set the appropriate `JAVA_HOME` path and the `FLUME_HOME` path in the `.bashrc` file.

To test if Flume has been installed correctly, we try executing `./flume-ng`. On successful installation, we get the following output.

```
dhruvsgarg@dhruvsgarg:~$ ls
anaconda3  Downloads  Hadoop      NLP_try.ipynb  Public      Template
Desktop    federated-averaging-tutorials-master  Music      nltk_data      summary.py   Videos
Documents  Flume      NLPPProject  Pictures        'TCS Internship'
```

```
dhruvsgarg@dhruvsgarg:~$ cd Flume/
dhruvsgarg@dhruvsgarg:~/Flume$ ls
bin          conf          doap_Flume.rdf  flume.conf  LICENSE  README.md  tools
CHANGELOG    DEVNOTES     docs            lib          NOTICE  RELEASE-NOTES
```

```
dhruvsgarg@dhruvsgarg:~/Flume$ cd bin/
dhruvsgarg@dhruvsgarg:~/Flume/bin$ ./flume-ng
Error: Unknown or unspecified command ''

Usage: ./flume-ng <command> [options]...
```

```
commands:
  help          display this help text
  agent         run a Flume agent
  avro-client   run an avro Flume client
  version       show Flume version info
```

```
global options:
  --conf,-c <conf>    use configs in <conf> directory
  --classpath,-C <cp> append to the classpath
  --dryrun,-d          do not actually start Flume, just print the command
  --plugins-path <dirs> colon-separated list of plugins.d directories. See the
```

STEP 2:

Setting up Hadoop version 3.2: We download and install **Hadoop 3.1.2**. We also set the appropriate JAVA_HOME path and HADOOP_HOME path along with YARN_HOME, etc. We must also manually add content to the various .xml files to set the appropriate configuration settings.

```
dhruvsgarg@dhruvsgarg:~$ hadoop version
Hadoop 3.1.2
Source code repository https://github.com/apache/hadoop.git -r 1019dde65bcf12e05ef48ac71e84550d589e5d9a
Compiled by sunilg on 2019-01-29T01:39Z
Compiled with protoc 2.5.0
From source with checksum 64b8bdd4ca6e77cce75a93eb09ab2a9
This command was run using /home/dhruvsgarg/Hadoop/share/hadoop/common/hadoop-common-3.1.2.jar
dhruvsgarg@dhruvsgarg:~$ cd $FLUME_HOME
dhruvsgarg@dhruvsgarg:~/Flume$ ls
bin  CHANGELOG  conf  DEVNOTES  doap_Flume.rdf  docs  lib  LICENSE  NOTICE  README.md  RELEASE-NOTES  tools
```

STEP 3:

Creating a Twitter application: After logging into Twitter, we create a new application using <https://apps.twitter.com>. Here I have created an application named **dhruvFirstFlumeApp**.

The screenshot shows the Twitter Developer app creation interface. On the left, there's a sidebar with links: 'Understanding apps', 'What is an app?', 'Why register an app?', and 'Which products require an API key?'. The main content area is titled 'App details' and contains the following information:

- App name (required):** A text input field containing 'dhruvFirstFlumeApp'. A note indicates 'Maximum characters: 32'.
- Application description (required):** A text area with the description 'To stream the Tweets using Apache Flume and to store them into HDFS.' A note indicates 'Share a description of your app. This description will be visible to users so this is a good place to tell them what your app does.'

STEP 4:

After creating the application, we can **access the keys and tokens** of the application.

There are 4 important application settings:

1. Consumer Key (API Key)
2. Consumer Secret (API Secret)
3. Access Token
4. Access Token Secret.

The screenshot shows the Twitter Developer console interface. At the top, there's a navigation bar with links: Developer, Use cases, Products, Docs, and More. Below this, the breadcrumb 'Apps > dhruvFirstFlumeApp' is visible. The main content area has three tabs: 'App details', 'Keys and tokens' (which is active), and 'Permissions'. Under the 'Keys and tokens' tab, there's a section titled 'Keys and tokens' with a subtitle 'Keys, secret keys and access tokens management.' Below this, there's a 'Consumer API keys' section showing an API key 'gmlV86TAcS3gJZGxBVXUkRnSM' and an API secret key 'x5GL2Igb4nofQZTe2DIW5RPN2Dtg1nm0jzblPSDMkL3VXVFmGL'. A 'Regenerate' button is present. The next section is 'Access token & access token secret', showing an access token '3096225511-015CCxecMIZOncG5dzRT8pj6brO5UJRXBdONh5I' and an access token secret 'eV9PSU4jcnunysaK1jC07aswBzDGk0x7GYEWXFXKpiOet'. The access level is 'Read and write'. There are 'Revoke' and 'Regenerate' buttons at the bottom.

STEP 5:

We create a **flume.conf** file in Flume's root directory as shown below. As discussed in the Flume Architecture, we configure our Source, Sink and Channel. In this application, the Source is Twitter, from where we are streaming our data and our Sink is HDFS, where we are writing the data.

The screenshot shows an IDE window with the 'flume.conf' file open. The menu bar includes File, Edit, View, Projects, Bookmarks, Sessions, Tools, Settings, and Help. The file content is as follows:

```

1 TwitterAgent.sources = Twitter
2 TwitterAgent.channels = MemChannel
3 TwitterAgent.sinks = HDFS
4
5 TwitterAgent.sources.Twitter.type = org.apache.flume.source.twitter.TwitterSource
6 TwitterAgent.sources.Twitter.channels = MemChannel
7 TwitterAgent.sources.Twitter.consumerKey = gmlV86TAcS3gJZGxBVXUkRnSM
8 TwitterAgent.sources.Twitter.consumerSecret = x5GL2Igb4nofQZTe2DIW5RPN2Dtg1nm0jzblPSDMkL3VXVFmGL
9 TwitterAgent.sources.Twitter.accessToken = 3096225511-015CCxecMIZOncG5dzRT8pj6brO5UJRXBdONh5I
10 TwitterAgent.sources.Twitter.accessTokenSecret = eV9PSU4jcnunysaK1jC07aswBzDGk0x7GYEWXFXKpiOet
11 TwitterAgent.sources.Twitter.keywords = hadoop, big data, analytics, data science, business intelligence, mapreduce
12

```

Source properties: In source configuration we are passing the Twitter source type as `org.apache.flume.source.twitter.TwitterSource`. Then, we are passing all the four tokens which we received from Twitter. At last in source configuration we are passing the keywords on which we are going to fetch the tweets.

```

13 TwitterAgent.sinks.HDFS.channel = MemChannel
14 TwitterAgent.sinks.HDFS.type = hdfs
15 TwitterAgent.sinks.HDFS.path = hdfs://localhost:9000/user/flume/tweets/
16 TwitterAgent.sinks.HDFS.fileType = DataStream
17 TwitterAgent.sinks.HDFS.writeFormat = Text
18 TwitterAgent.sinks.HDFS.batchSize = 1000
19 TwitterAgent.sinks.HDFS.rollSize = 0
20 TwitterAgent.sinks.HDFS.rollCount = 10000
21 TwitterAgent.sinks.HDFS.rollInterval = 600

```

Sink properties: In the sink configuration, we set the type, path, file type, batch size and roll parameters.

```

22
23 TwitterAgent.channels.MemChannel.type = memory
24 TwitterAgent.channels.MemChannel.capacity = 10000
25 TwitterAgent.channels.MemChannel.transactionCapacity = 100
26

```

Channel properties: In the channel properties, we set the memory, capacity and the capacity of each transaction.

STEP 6:

After setting up all the required tools, we are **set for execution**. We perform the Twitter Data Streaming execution using the following command:

```
$FLUME_HOME/bin/flume-ng agent --conf ./conf/ -f $FLUME_HOME/flume.conf
```

```

dhruvsgarg@dhruvsgarg:~/Flume$ $FLUME_HOME/bin/flume-ng agent --conf ./conf/ -f $FLUME_HOME/flume.conf
Info: Sourcing environment configuration script /home/dhruvsgarg/Flume/conf/flume-env.sh
Info: Including Hadoop libraries found via (/home/dhruvsgarg/Hadoop/bin/hadoop) for HDFS access

```

We stop the execution by entering Ctrl+C after some time.

STEP 7:

After some time, the **execution completes** and the data has been streamed into HDFS from Twitter.

```

Firmly believe every South African should have a share trading account. Let's get involved:
bulali@mp44.co.za(2017-11-15T00:06:21Z)Please leave black kings out if this mess https://
t.co/UbLLLnzhER(2017-11-15T00:06:21Z)Twitter for Android<a href="http://twitter.com/download/android" rel="nofollow">Twitter for Android</a>
$93050469406994568(2017-11-15T00:06:21Z)Gordon/Barrie Island, OntarioJesus follower. husband. dad. cow-calf farmer on Manitoulin Island with
@Grandview_Farms(2017-11-15T00:06:21Z)Barstow (Jordan)barstowmiller(2017-11-15T00:06:21Z)OntarioBeef cattle moving season isn't hard to
summarize in pics & always best when synonymous with #JohnnyCash se... https://t.co/r7XhSIUjUz(2017-11-15T00:06:21Z)This page is for all #music lovers . Like us
& join our group where you can post & #enjoy your favorite #music #FENDER #GIBSON #MARSHALL #TEAMFOLLOWBACK(2017-11-15T00:06:21Z)LOVE OF MUSIC
LOVE_OF_MUSIC(2017-11-15T00:06:21Z)BrunoMars Just the way you Are #Music https://t.co/eCCe0DKEXn(2017-11-15T00:06:21Z)The Social Jukebox<a href="https://www.socialjukebox.com" rel="nofollow">The Social Jukebox</a>
$930504698259312640(2017-11-15T00:06:21Z)テイニストロード次元ネファリア東大生
と即ハメしたりマジックしたりモンハンしたりするbotですけどbotじゃないです (半ギレ)

・音ゲーマ〜見習い (たのしかったダンエボ 武器: 2554-8213)
・MtGスタンモダンコマンドリーミテMO(2017-11-15T00:06:21Z)nobikybs_bot(2017-11-15T00:06:21Z)ハッモしかしてツイッター
社の中に王宮の弾圧の使い手が... ?(2017-11-15T00:06:21Z)twicca(2017-11-15T00:06:21Z)RT @usernameimkey: This is so important.
Check in, it could save someone's life. https://t.co/PAYlgPbjGe(2017-11-15T00:06:21Z)Twitter for iPhone<a href="http://twitter.com/download/iphone"
rel="nofollow">Twitter for iPhone</a>
$930504698238443520(2017-11-15T00:06:21Z)Bps memory (2017-11-15T00:06:21Z)ARYKAWAAI

```

RESULT

Thus Flume was successfully used for Twitter Data Streaming and the output was stored in HDFS.

OTHER APPLICATIONS OF APACHE FLUME

- Flume is used in e-commerce companies to analyse the customer behaviour of different regions.
- It is used to feed huge log data generated by application servers into HDFS at a higher speed.
- Flume is used in online analytics application hence it is a backbone for real-time event processing.
- Flume efficiently feed log data files from multiple web servers into a centralized store i.e HDFS or HBase.
- Flume can process data in-flight using interceptors. These can be very useful for data masking or filtering.
- Flume is distributed in nature i.e agents can be installed on many machine.
- Flume is able to collect real-time and batch mode data from multiple servers and immediately feed into Hadoop.
- Flume is also used to import huge volumes of event data (instantaneous values) produced by social networking sites like Facebook and Twitter, and e-commerce websites like Amazon and Flipkart.